

Introduction to Robotics

Search and Motion Planning

WS 2015 / 16

Prof. Dr. Daniel Göhring
Intelligent Systems and Robotics
Department of Computer Science
Freie Universität Berlin

State-Space Search

- States Z (Vertices)
- Operators $Op \subseteq Z \times Z$ (Edges)
- Initial states $z_{\text{initial}} \in Z$
- Target states $Z_{\text{final}} \subseteq Z$

Graph $[Z, Op]$

- Cost functions $c: Z \times Z \rightarrow R^+$

Costs of a path $w = z_0 z_1 \dots z_n \in Z^*$:

$$c(z_0 z_1 \dots z_n) := \sum_{i=1, \dots, n} c(z_{i-1}, z_i)$$

- Estimator function $\sigma: Z \rightarrow R$ (Heuristic) for remaining costs from z to a target state.

Slides: H.-D. Burkhard

State-Space Search

- Tasks:
- Can a goal state $z \in Z_{\text{final}}$ be reached from initial state $z_{\text{initial}} \in Z$?
- Find a way from initial state $z_{\text{initial}} \in Z$ to a target state $z \in Z_{\text{final}}$
- Find an optimal path from initial state $z_{\text{initial}} \in Z$ to a target state $z \in Z_{\text{final}}$

Complexity (Number of States / Vertices)

- 8 Puzzle: $9!$ states
 $9!/2 = 181.440$ reachable
- 15 Puzzle: $16!$ states
 $16!/2$ reachable
- Rubik's cube: $12 \cdot 4,3 \cdot 10^{19}$ states
 $1/12$ reachable: $4,3 \cdot 10^{19}$
- Towers of Hanoi: 3^n States for n discs
solvable in $(2^n) - 1$ moves
- Checkers: approx. 10^{40} games of average length
- Chess: approx. 10^{120} games of average length
- Go: 3^{361} states

Expansion Strategies

– Directions

- Forward, start with z_{initial}
(forward chaining, data driven, bottom up)
- Reverse z_{final}
(backward chaining, goal driven, top down)
- Bi-directional

– Expansion


- Depth first
- Breadth first

– Additional information

- blind search („uninformed“)
- heuristisch search with σ („informed“)

Bi-Direktional Breadth-Search

- Parallel search from start and goal until meeting



- Search depth from both sides only half

Expansion

Data structures :

- OPEN List:

A vertex is "open" , if it was constructed but not expanded (neighboring vertices not calculated)

- CLOSED List:

A vertex is „closed" , if it was fully expanded (all neighboring vertices are known)

Further information:

Predecessor / successor of vertices
for reconstruction of found paths

Heuristic Search for best Way: A*

Costs to reach z' from z :

– If z' is reachable from z :

$$g(z, z') := \text{Min} \{ c(s) / s \text{ path from } z \text{ to } z' \},$$

– Else: $g(z, z') := \infty$

Tentative cost calculation during expansion:


$G' = [V', E']$ is a (known) partial graph of G

$$g'(z, z', G') := \text{Min} \{ c(s) / s \text{ path in } G' \text{ from } z \text{ to } z' \}$$

$$g'(z, z', G') \geq g(z, z')$$

Heuristic Search for Best Path

-



$g'(z_0, z', G')$: so far known costs to reach z' from Start

$\sigma(z')$: estimated costs to reach target state, starting from z'

Algorithm A* („soft form“)

For Trees

- A*0: (Start) $OPEN := [z_0]$, $CLOSED := []$.
- A*1: (negative exit)
 If $OPEN = []$: EXIT(„no“).
- A*2: (positive exit)
 If z first vertice in $OPEN$:
 If z is target: EXIT(„yes:“ z).
- A*3: (expand)
 $OPEN := OPEN - \{z\}$. $CLOSED := CLOSED \cup \{z\}$.
 $Succ(z) :=$ set of successors of z .
 If $Succ(z) = \{\}$: Goto A*1.
- A*4: (Organization of $OPEN$)
 – $OPEN := OPEN \cup Succ(z)$ with increasing order of $g'(z_0, z', G') + \sigma(z')$
- Goto A*1.

Search space as a tree:
CLOSED not used.

Algorithm A* („soft form“)

For Trees and for Cyclic Graphs

- A*0: (Start) $OPEN := [z_0]$, $CLOSED := []$.
- A*1: (negative exit)
 If $OPEN = []$: EXIT(„no“).
- A*2: (positive exit)
 If z first vertice in $OPEN$:
 If z is target: EXIT(„yes:“ z).
- A*3: (expand)
 $OPEN := OPEN - \{z\}$. $CLOSED := CLOSED \cup \{z\}$.
 $Succ(z) :=$ set of successors of z .
 If $Succ(z) = \{\}$: Goto A*1.
 $NEW = Succ(z) - \{z' \mid z' \in Succ(z) \text{ and } z' \in CLOSED \text{ and } g'(z_0, z', G') \geq g'(z_0, z', G'_{old})\}$
- A*4: (Organization of $OPEN$)
 - $OPEN := OPEN \cup NEW$ with increasing order of $g'(z_0, z', G') + \sigma(z')$
 - If elements z' occur twice in $OPEN$, delete entry with bigger costs.
- Goto A*1.

If search space has cycles:
use **CLOSED**.

Algorithm A* („soft form“)

Definition: $f(z) := \text{Min} \{ g(z, z_{\text{final}}) \mid z_{\text{final}} \in Z_{\text{final}} \}$

= real costs from z to target state (vertex)

($f(z_0)$ = cost of the optimal path)

Heuristic function σ is called optimistic (a.k.a. admissible) or underestimating,

if $\sigma(z) \leq f(z)$ for all $z \in Z$.

Proposition :

Given.: Ex. $\delta > 0$ with $c(z, z') > \delta$ for all z, z' .

σ is optimistic heuristic.

every node has a finite number of successors.

One can prove: If solution exists, A* (soft form) finds an optimal path.

Special Cases

$c \equiv 0$:

Search for best path with heuristic σ and no cost function
(Hill climbing)

$\sigma \equiv 0$:

Search for best path without heuristic
($\sigma \equiv 0$ ist also optimistic heuristic) - Dijkstra

$c \equiv 1$ ($g' \equiv \text{Search depth}$) , $\sigma \equiv 0$:

Breadth first search

Influence of Heuristic σ


- Trade-off between quality of result and calculation effort
- Criteria for order of extension
- Same order as $g' + \sigma$ is also provided by $a \cdot (g' + \sigma) + b$ for arbitrary positive constants a, b .
- optimal order for $\sigma = f$
- σ_2 more effective for σ_1 if $\sigma_1 \leq \sigma_2 \leq f$
(Hierarchy for heuristic functions)

Algorithm A* („hard form“)

- A*0: (Start) $OPEN := [z_0]$, $CLOSED := []$.
- A*1: (negative exit)
 Falls $OPEN = []$: EXIT(„no“).
- A*2: (positive exit)
 If z first state in $OPEN$.
 If z target state: EXIT(„yes:“ z).
- A*3: (expand)
 $OPEN := OPEN - \{z\}$. $CLOSED := CLOSED \cup \{z\}$.
 $Succ(z) :=$ set of successors of z .
 If $Succ(z) = \{\}$: Goto A*1.
- A*4: (Organize of $OPEN$)
 $OPEN := OPEN \cup (Succ(z) - CLOSED)$ sorted by increasing
 $g'(z_0, z', G') + \sigma(z')$
- Goto A*1.

Algorithmus A* („hard form“)

- Problem:
- For optimistic σ the hard form is not always correct



Algorithm A* („hard form“)

Definition:

Heuristic σ is called consistent,
if for all states z', z'' holds:

$$\sigma(z') \leq g(z', z'') + \sigma(z'')$$

Lemma: If σ is consistent, σ is optimistic.

(The reverse does not necessarily apply)

Algorithm A* („hard form“)

Proposition :

Given: Ex. $\delta > 0$ with $c(z, z') > \delta$ for all z, z' .

σ is a consistent heuristic

One can prove: if a solution exists, A* (hard form) finds an optimal path

Algorithm A* („hard form“)

- „consistent“ functions harder to find than „optimistic“
- It is also possible to use an optimistic heuristic with weaker pruning of search space
- Erase states in OPEN (or Succ(z)) only if new evaluation is worse than earlier evaluation.

Algorithm A*

- Search costs ~ Number of states to expand,
- ~ Computational effort of σ and g'
- Space requirements exponential
- Search costs vs. costs of solution (optimal / suboptimal solutions)
- Measure: Path length/expanded vertices

Memory Saving Variants of Algorithm A*


- Iterative Deepening A* (IDA*) in analogy to IDA:
 - Depth first till boundary $g'(z_0, z', G') + \sigma(z')$ is exceeded from earlier iteration
- SMA* (simplified memory-bounded A*)

Anytime (or Weighted) A*

- If the heuristic σ (sometimes also denoted h) is closer to the real costs, less vertices have to be expanded
- But „inflating“ the heuristics can lead to a heuristic which is not optimistic (admissible),
(the following slides with courtesy to Likhachev)

Heuristics in Heuristic Search

- Dijkstra's: expands states in the order of $f = g$ values
- A* Search: expands states in the order of $f = g + h$ values
- **Weighted A***: expands states in the order of $f = g + \epsilon h$ values, $\epsilon > 1$ = bias towards states that are closer to goal




Suboptimality


- Heuristic includes factor ϵ
- Suboptimality is bounded by factor ϵ
 - The length of the found solution is not longer than ϵ times the optimal solution
- Exampe:
 - costs from cell to cell are 1
 - Heuristic is the larger of coordinate difference from current cell to goal cell
 - Start is upper left cell
 - Goal is lower right
 - Obstacles black
 - Free space white
 - Expanded cells gray

Anytime Search based on weighted A*


- Constructing anytime search based on weighted A*:
 - find the best path possible given some amount of time for planning
 - do it by running a series of weighted A* searches with decreasing ϵ :



13 expansions
solution=11 moves



15 expansions
solution=11 moves



20 expansions
solution=10 moves

Anytime A*

- Problem:
 - Running A* with increasing ϵ each time from scratch can be very expensive
 - Many states remain the same through various iterations
 - A solution for reusing the search results is described in ARA* (Likhachev)
- ARA* : an efficient version of the above that reuses state values within any search iteration

ARA*

- Efficient series of weighted A* searches with decreasing ϵ :

ComputePathwithReuse function

```
while( $f(s_{goal}) > \text{minimum } f\text{-value in OPEN}$ )  
  remove  $s$  with the smallest  $[g(s) + \epsilon h(s)]$  from OPEN;  
  insert  $s$  into CLOSED;  
  for every successor  $s'$  of  $s$   
    if  $g(s') > g(s) + c(s, s')$   
       $g(s') = g(s) + c(s, s')$ ;  
    if  $s'$  not in CLOSED then insert  $s'$  into OPEN;  
    otherwise insert  $s'$  into INCONS
```

ARA*

- Efficient series of weighted A* searches with decreasing ϵ :

ComputePathwithReuse function


```
while( $f(s_{goal}) > \text{minimum } f\text{-value in OPEN}$ )  
  remove  $s$  with the smallest  $[g(s) + \epsilon h(s)]$  from OPEN;  
  insert  $s$  into CLOSED;  
  for every successor  $s'$  of  $s$   
    if  $g(s') > g(s) + c(s, s')$   
       $g(s') = g(s) + c(s, s')$ ;  
    if  $s'$  not in CLOSED then insert  $s'$  into OPEN;  
    otherwise insert  $s'$  into INCONS
```

```
set  $\epsilon$  to large value;  
 $g(s_{start}) = 0$ ; OPEN =  $\{s_{start}\}$ ;  
while  $\epsilon \geq 1$   
  CLOSED =  $\{\}$ ; INCONS =  $\{\}$ ;  
  ComputePathwithReuse();  
  publish current  $\epsilon$  suboptimal solution;  
  decrease  $\epsilon$ ;  
  initialize OPEN = OPEN  $\cup$  INCONS;
```

ARA*


- A series of weighted A* searches

$\epsilon = 2.5$




13 expansions
solution=11 moves

$\epsilon = 1.5$



15 expansions
solution=11 moves


$\epsilon = 1.0$



20 expansions
solution=10 moves


- ARA*

$\epsilon = 2.5$




13 expansions
solution=11 moves

$\epsilon = 1.5$



1 expansion
solution=11 moves


$\epsilon = 1.0$



9 expansions
solution=10 moves

D* and Variations

- D* and its variations have been used on Mars rovers Opportunity and Spirit and by CMU at the DARPA Grand Challenge
- Roughly:
 - D* works as A* but from goal to target
 - Every expanded node „knows“ its predecessor
 - When start node (vertex) is the next node, the search is done
 - modes are marked: NEW (was never in OPEN), OPEN, CLOSE (no longer in OPEN), LOWER, RAISE (cost is higher than last time in OPEN)



Literature

- J.C. Latombe, Robot Motion Planning, Kluwer Academic Publishers, 1991.
- Hans-Dieter Burkhard: Einführung in die Künstliche Intelligenz (2008)
- Likhachev, ARA*, CMU
- Choset, Motion Planning, CMU,
- Toussaint, Lecture Notes Robotics, 2011
- Dr. John (Jizhong) Xiao, City College New York