

Introduction to Robotics

Motion Planning II


WS 2015 / 16

Prof. Dr. Daniel Göhring
Intelligent Systems and Robotics
Department of Computer Science
Freie Universität Berlin

Heuristics in Heuristic Search

- Dijkstra's: expands states in the order of $f = g$ values

What are the states expanded?




*Slides: Maxim
Likhachev*

Heuristics in Heuristic Search


- A* Search: expands states in the order of $f = g + h$ values

for high-D problems, this results in A being too slow and running out of memory*



Heuristics in Heuristic Search


- Weighted A* Search: expands states in the order of $f = g + \varepsilon h$ values, $\varepsilon > 1$
- bias towards states that are closer to goal



Robot Motion Planning

- Search in geometric structures
- ***Spatial reasoning***
- Challenges:
 - Continuous state space
 - Large dimensional space
- Application of other search approaches like A*

Degrees of Freedom - Revisited




- The geometric configuration of a robot is defined by p degrees of freedom (DOF)
- Assuming p DOFs, the geometric configuration A of a robot is defined by p variables:


$A(\mathbf{q})$ with $\mathbf{q} = (q_1, \dots, q_p)$

- Examples:
 - Prismatic (translational) DOF: q_i is the amount of translation in some direction
 - Rotational DOF: q_i is the amount of rotation about some axis

Examples




Allowed to move only
in x and y : 2DOF




Allowed to move in x
and y and to rotate:
3DOF (x, y, θ)


Examples



Fixed (attached at the base)




Free Flying




Fixed (the dashed line is constrained to be horizontal)

Configuration (C-Space)




$$\mathbf{q} = (x, y, \theta)$$
$$\mathcal{C} = \mathbb{R}^2 \times \text{set of 2-D rotations}$$




$$\mathbf{q} = (q_1, q_2)$$
$$\mathcal{C} = \text{2-D rotations} \times \text{2-D rotations}$$

- Configuration space \mathcal{C} = set of values of \mathbf{q} corresponding to legal configurations of the robot
- Defines the set of possible parameters (the search space) and the set of allowed paths

Configuration Space



2R manipulator



Configuration space

Motion Planning


- The world consists of
 - Obstacles:
 - already occupied space,
 - robot cannot go there
 - Free Space
 - Space which is not occupied
 - Robot „might“ go there
 - Determine if it can go there
 - Start position and goal position

- In a nutshell:

Find a collision-free trajectory under physical, geometrical and temporal constraints from start to goal.

World Frame and Robot Frame


- A: single rigid object –(the robot)
- W: Euclidean space where A moves;
- $W = \mathbb{R}^2$ or \mathbb{R}^3
- B_1, \dots, B_m : fixed rigid obstacles distributed in W



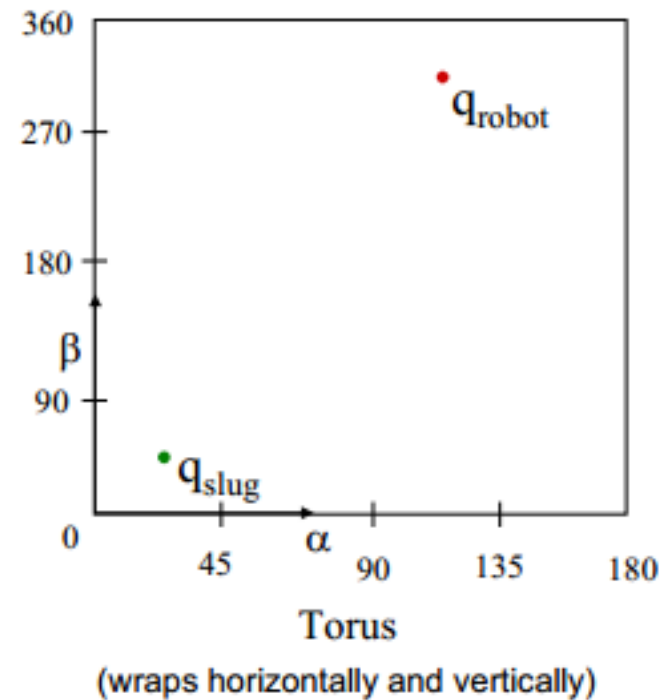
F_W : World frame
 F_A : Robot frame, a moving frame rigidly associated with the robot

- Configuration q of A is a specification of the position and orientation of A in world frame F_W .
- Configuration space is the space of all possible robot configurations

Configuration Space




Two points in the robot's workspace



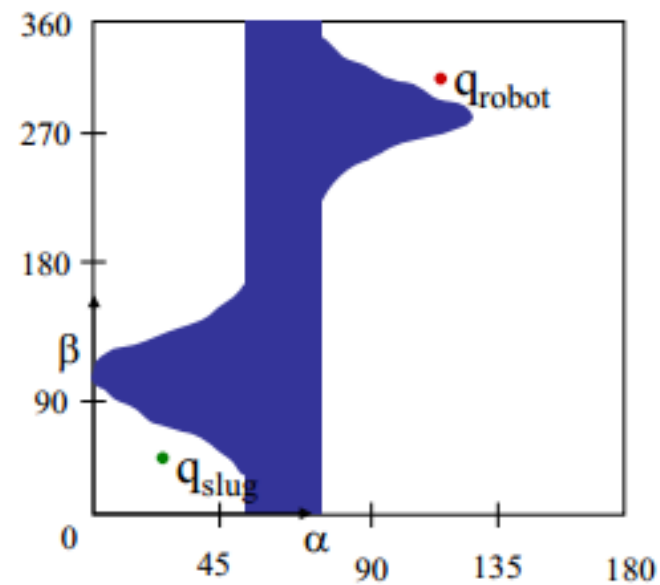
Objects in Configuration (or C)-Space

If the robot configuration is within the blue area, it will hit the obstacle



An obstacle in the robot's workspace


What is dimension of the C-space of puma robot (6R)?



a "C-space" representation

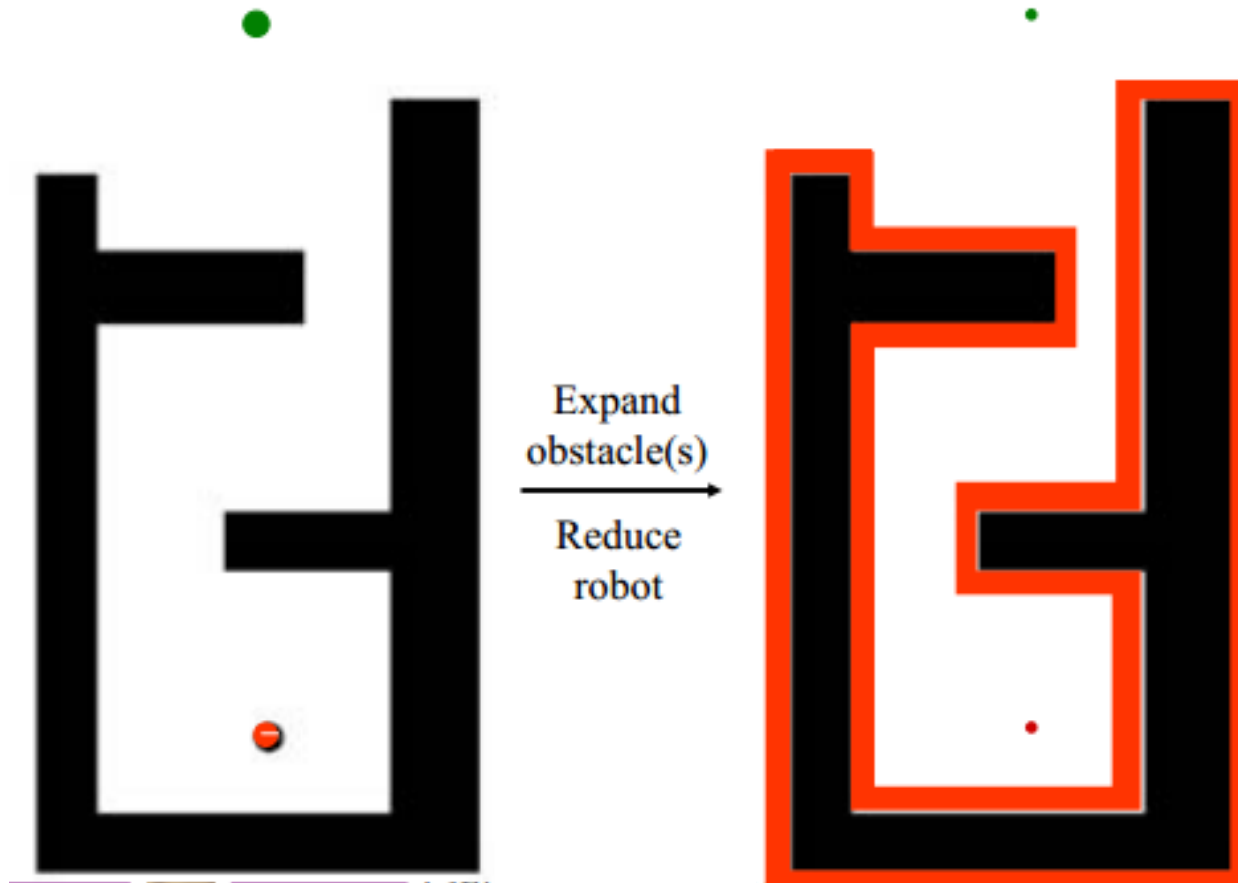
Visualization of high dimension C-space is difficult

Free Space, Point Robot




- $\mathcal{C}_{\text{free}} = \{\text{Set of parameters } \mathbf{q} \text{ for which } A(\mathbf{q}) \text{ does not intersect obstacles}\}$
- For a point robot in the 2-D plane: \mathbb{R}^2 minus the obstacle regions


What if the Robot is not a Point



Free Space: Symmetric Robot




Free Space: Symmetric Robot



- We still have $\mathcal{C} = \mathbb{R}^2$ because orientation does not matter
- Reduce the problem to a point robot by expanding the obstacles by the radius of the robot


Free Space: Non-Symmetric Robot



- The configuration space is now three-dimensional (x, y, θ)
- We need to apply a different obstacle expansion for each value of θ
- We still reduce the problem to a point robot by expanding the obstacles

Free Space

C-OBSTACLE REGION




FREE SPACE

$$C_{\text{free}} = C \setminus \bigcup_{i=1}^m CB_i = \{q \in C : A(q) \cap \bigcup_{i=1}^m CB_i = \emptyset\}$$

Free configuration q iff $q \in C_{\text{free}}$

From
Robot Motion Planning
J.C. Latombe


More Complex Configuration-Spaces



- In all cases: The problem is reduced to finding the path of a *point* through configuration space by “expanding the obstacles”

Minkowski Sum, Generating Obstacles in C-Space

This expansion of one planar shape by another is called the *Minkowski sum* \oplus



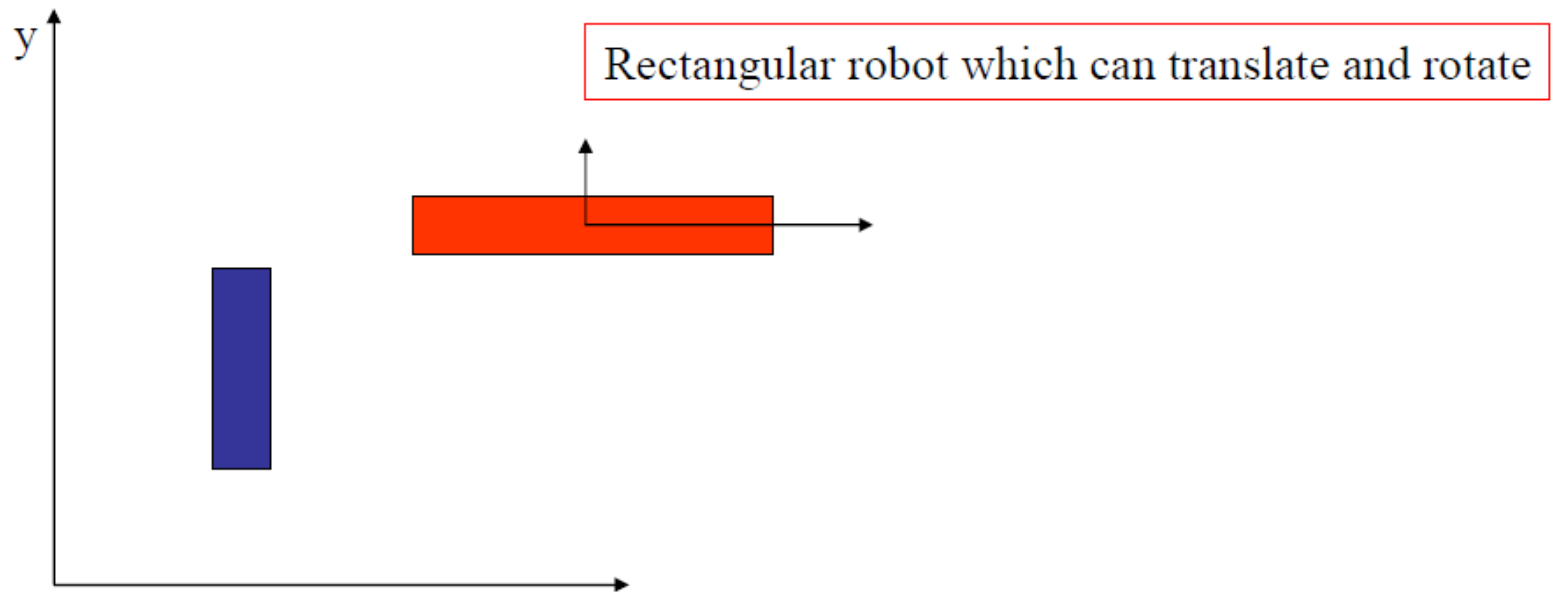
(Dilation operation)
$$P \oplus R = \{ p + r \mid p \in P \text{ and } r \in R \}$$

Used in robotics to ensure that there are free paths available.

Additional Dimension

What would the \mathcal{C} -obstacle be if the rectangular robot (red) can *translate and rotate* in the plane.


(The blue rectangle is an obstacle.)




C-Obstacle in 3D

What would the C -obstacle be if the rectangular robot (red) can *translate and rotate* in the plane.

(The blue rectangle is an obstacle.)




3-D
configuration space




C-Obstacle in 3D

What would the configuration space of a 3DOF rectangular robot (red) in this world look like?

(The obstacle is blue.)





3-D
configuration space





One Slice

Taking one slice of the \mathcal{C} -obstacle in which the robot is rotated 45 degrees...




2D-Projection





why not keep it this simple?

Projection Problems



Motion Planning Problem



- A = robot with p degrees of freedom in 2-D or 3-D
- CB = Set of obstacles
- A configuration \mathbf{q} is legal if it does not cause the robot to intersect the obstacles
- Given start and goal configurations (\mathbf{q}_{start} and \mathbf{q}_{goal}), find a continuous sequence of legal configurations from \mathbf{q}_{start} to \mathbf{q}_{goal} .
- Report failure if not path is found

Topics

- Visibility Graphs
 - Roadmaps
 - Voronoi
- Approximate Cell Decomposition
- Potential Fields
- Probabilistic Roadmaps (Sampling)


Topics

- Visibility Graphs
 - Roadmaps
 - Voronoi
- Approximate Cell Decomposition
- Potential Fields
- Probabilistic Roadmaps (Sampling)

In all cases:


- Reduce the intractable problem in continuous C-space to a tractable problem in a discrete space Use all of the techniques we know (A^* , stochastic search, etc.)

Visibility Graphs



In the absence of obstacles, the best path is the straight line between q_{start} and q_{goal}

Visibility Graphs




- Assuming polygonal obstacles: It looks like the shortest path is a sequence of straight lines joining the vertices of the obstacles.
- Is this always true?

Roadmap Approaches, Visibility Graphs

- used to find Euclidean shortest paths among a set of polygonal obstacles in the plane
- construct all of the line segments that connect vertices to one another (and that do not intersect the obstacles themselves)
- Formed by connecting all “visible” vertices, the start point and the end point, to each other
- For two points to be “visible” no obstacle can exist between them
- Paths exist on the perimeter of obstacles
- a graph is defined, converts the problem into graph search.
- Dijkstra’s algorithm, $O(N^2)$, N = the number of vertices in C-space


The Visibility Graph in Action (1)

- First, draw lines of sight from the start and goal to all “visible” vertices and corners of the world.




The Visibility Graph in Action (2)

- Second, draw lines of sight from every vertex of every obstacle like before. Remember lines along edges are also lines of sight.




The Visibility Graph in Action (3)

- Repeat until done



Optimality

- Visibility graphs do not preserve their optimality in higher dimensions




- In addition, the paths they find are “semi-free,” i.e. in contact with obstacles.
- No clearance, any execution error will lead to a collision
- Need other types of roadmaps for safer paths

Voronoi Diagrams




- Given a set of data points in the plane:
 - Color the entire plane such that the color of any point in the plane is the same as the color of its nearest neighbor

Voronoi Diagrams




- Voronoi diagram = The set of line segments separating the regions corresponding to different colors
 - Line segment = points equidistant from 2 data points
 - Vertices = points equidistant from > 2 data points

Voronoi Diagrams




- Complexity (in the plane):
- $O(N \log N)$ time
- $O(N)$ space

Voronoi Diagram




Voronoi Diagrams: Beyond Points




- Edges are combinations of straight line segments and segments of quadratic curves
- Straight edges: Points equidistant from 2 lines
- Curved edges: Points equidistant from one corner and one line

Voronoi Diagrams (Polygons)




Voronoi Diagrams: Metric

- Many ways to measure distance; two are:
 - L_1 metric
 - $(x,y) : |x| + |y| = \text{const}$
 - L_2 metric
 - $(x,y) : x^2 + y^2 = \text{const}$



$$\{(x,y) : |x| + |y| = \text{const}\}$$

L1




$$\{(x,y) : x^2 + y^2 = \text{const}\}$$

L2

Minkowski Metric


$$\left(\sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

Source: Wikipedia



Voronoi Diagram: Construction (Roughly)


- For a Voronoi-Diagram in 2D: there are many options, e.g.: Fortune's Algorithm (a form of Sweepline Algorithm)
- Using $O(n \log n)$ time and
- $O(n)$ space




- Find intersections of 2 parabolas (breakpoints) and intersection points of three parabolas while sweeping down


Voronoi Diagram: Construction (Roughly)

- For a Voronoi-Diagram in 2D: there are many options, e.g.: Fortune's Algorithm (a form of Sweepline Algorithm)
- Using $O(n \log n)$ time and
- $O(n)$ space






Voronoi Diagram (Polygons)



- Key property: The points on the edges of the Voronoi diagram are the *furthest from the obstacles*
- Idea: Construct a path between q_{start} and q_{goal} by following edges on the Voronoi diagram
- (Use the Voronoi diagram as a roadmap graph instead of the visibility graph)

Voronoi-Diagram (Planning)



- Find the point q^*_{start} **of the Voronoi** diagram closest to q_{start}
- Find the point q^*_{goal} **of the Voronoi** diagram closest to q_{goal}
- Compute shortest path from q^*_{start} **to q^*_{goal} on the Voronoi diagram**

Voronoi: Weaknesses

- Difficult to compute in higher dimensions or nonpolygonal worlds
- Approximate algorithms exist
- Use of Voronoi is not necessarily the best heuristic (“stay away from obstacles”) can lead to paths that are much too conservative
- Can be unstable → Small changes in obstacle configuration can lead to large changes in the diagram


Topics

- Visibility Graphs
 - Roadmaps
 - Voronoi
- Approximate Cell Decomposition
- Potential Fields
- Probabilistic Roadmaps (Sampling)
- Bug algorithms




Decompose the space into cells so that any path inside a cell is obstacle free

Approximate Cell Decomposition




- Define a discrete grid in C-Space
- Mark any cell of the grid that intersects C_{obs} as blocked
- Find path through remaining cells by using (for example) A^* (e.g., use Euclidean distance as heuristic)
- Cannot be *complete as described so far. Why?*

Approximate Cell Decomposition





Approximate Cell Decomposition



- Cannot find a path in this case even though one exists
- Solution:
- Distinguish between
 - Cells that are entirely contained in C_{obs} (*FULL*) and
 - Cells that partially intersect C_{obs} (*MIXED*)
- Try to find a path using the current set of cells
- If no path found:
 - Subdivide the *MIXED* cells and try again with the new set of cells

Quadtree Decomposition




 empty


 mixed

 full


Octree Decomposition



Example




Examples



Approximate Cell Decomposition: Limitations

- Good:
 - Limited assumptions on obstacle configuration
 - Approach used in practice
 - Find obvious solutions quickly
- Bad:
 - No clear notion of optimality (“best” path)
 - Trade-off completeness/computation
 - Still difficult to use in high dimensions

Exact Cell Decomposition




Any path within one cell is guaranteed to not intersect any obstacle

The diagram illustrates a 2D environment with a trapezoidal boundary. Inside, there are three black polygonal obstacles. Vertical red lines represent barriers or discretized steps. A blue path starts at a yellow circle labeled q_{start} and ends at a yellow circle labeled q_{end} . The path is composed of segments connecting various square markers: yellow squares are located on the path segments, and light blue squares are located at the intersections of the path with the vertical red lines. The path successfully navigates around the obstacles and through the barriers.

- Prof. Dr. Daniel Göhring

Exact Cell Decomposition




- Critical event: Create new cell
- Critical event: Split cell

Plane Sweep Algorithm


- Initialize current list of cells to empty
- Order the vertices of C_{obs} along the *x direction*
- For every vertex:
 - Construct the plane at the corresponding *x location*
 - Depending on the type of event:
 - Split a current cell into 2 new cells OR
 - Merge two of the current cells
 - Create a new cell

Exact Cell Decomposition



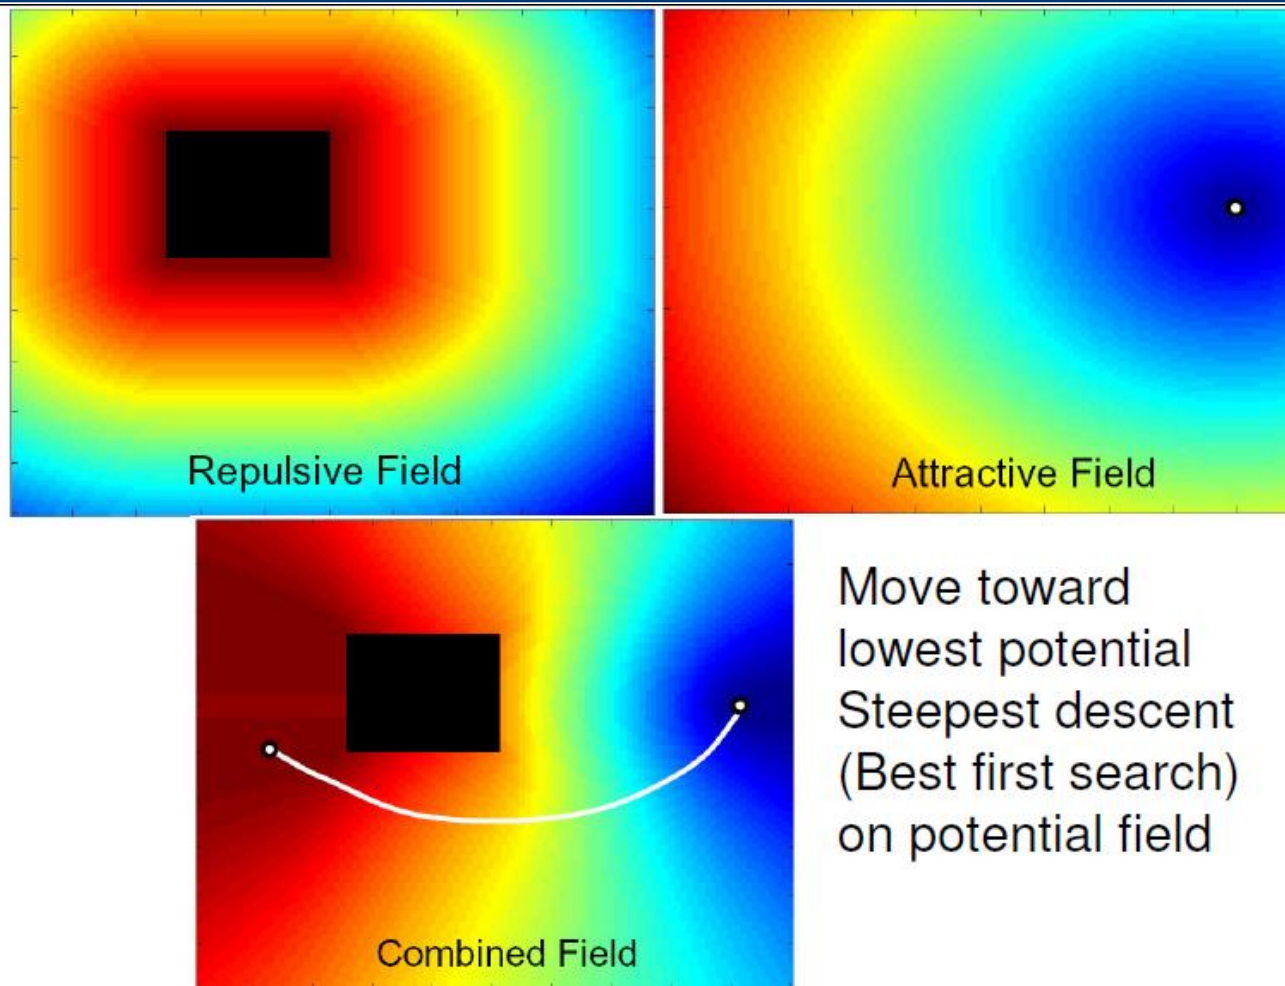
- A version of exact cell decomposition can be extended to higher dimensions and non-polygonal boundaries (“cylindrical cell decomposition”)
- Provides exact solution completeness
- Expensive and difficult to implement in higher dimensions

Potential Fields



- Stay away from obstacles: Imagine that the obstacles are made of a material that generate a *repulsive field*
- Move closer to the goal: Imagine that the goal location is a particle that generates an *attractive* field

Potential Fields



Representation Potential Fields

$$U_g(\mathbf{q}) = d^2(\mathbf{q}, \mathbf{q}_{goal})$$

Distance to goal state

$$U_o(\mathbf{q}) = \frac{1}{d^2(\mathbf{q}, Obstacles)}$$





Distance to nearest obstacle point.
Note: Can be computed efficiently by
using the *distance transform*

$$U(\mathbf{q}) = U_g(\mathbf{q}) + \lambda U_o(\mathbf{q})$$

λ controls how far we
stay from the obstacles

Potential Fields, Representation


- Closed Form as a math. formular,
- Gridbased




Potential field method

- After the total potential is obtained, generate force field (negative gradient)
- Let the sum of the forces control the robot.


Equipotential contours



Negative gradient




Total potential




To a large extent, this is computable from sensor readings.

Potential Fields: Limitations



- Completeness
- Problems in higher dimensions

Local Minimum Problem



- Potential fields in general exhibit local minima
- Special case: Navigation function
 - $U(q_{goal}) = 0$
 - For any q **different from** q_{goal} , **there exists a** neighbor q' **such that** $U(q') < U(q)$

Getting out of Local Minima I

- Repeat
 - If $U(q) = 0$ return **Success**
 - If too many iterations return **Failure**
 - Else:
 - Find neighbor q_n of q with **smallest $U(q_n)$**
 - If $U(q_n) < U(q)$ **OR q_n has not yet been** visited
 - Move to q_n ($q \leftarrow q_n$)
 - Remember q_n

May take a long time to explore region “around” local minima

Getting out of Local Minima II

- Repeat
 - If $U(\mathbf{q}) = 0$ return **Success**
 - If too many iterations return **Failure**
 - Else:
 - Find neighbor \mathbf{q}_n of \mathbf{q} with **smallest** $U(\mathbf{q}_n)$
 - If $U(\mathbf{q}_n) < U(\mathbf{q})$
 - Move to \mathbf{q}_n ($\mathbf{q} \leftarrow \mathbf{q}_n$)
 - Else
 - Take a random walk for T steps starting at \mathbf{q}_n
 - Set \mathbf{q} to the **configuration reached at the end of the random walk**


Topics

- Visibility Graphs
 - Roadmaps
 - Voronoi
- Approximate Cell Decomposition
- Potential Fields
- Probabilistic Roadmaps (Sampling)

Completely describing and optimally exploring the C-space is too hard in high dimension + it is not necessary → Limit ourselves to finding a “good” sampling of the C-space




Sampling Techniques (Probabilistic Roadmaps)




Sampling Techniques

Sample random locations




Sampling Techniques

Remove the samples in the forbidden regions




Sampling Techniques

Link each sample to its K nearest neighbors




Sampling Techniques

Remove the links that cross forbidden regions



Sampling Techniques


Remove the links that cross forbidden regions



The resulting graph is a *probabilistic roadmap (PRM)*

Sampling Techniques

Link the start and goal to the PRM and search using A*



Sampling Techniques

Continuous Space




Discretization



A* Search


- “Good” sampling strategies are important:
 - Uniform sampling
 - Sample more near points with few neighbors
 - Sample more close to the obstacles
 - Use pre-computed sequence of samples



Sampling Techniques

- Remarkably, we can find a solution by using *relatively few randomly* sampled points.
- In most problems, a relatively small number of samples is sufficient to cover most of the feasible space with probability 1
- For a large class of problems:
 - $\text{Prob}(\text{finding a path}) \rightarrow 1$ exponentially with the number of samples
- *But*, cannot detect that a path does not exist


Sampling Techniques (Probabilistic Roadmaps), Contd.



PRMs: The Good News

1. PRMs are probabilistically complete
2. PRMs apply easily to high-dimensional C-space
3. PRMs support fast queries w/ enough preprocessing

Many success stories where PRMs solve previously unsolved problems



PRMs: The Bad News

1. PRMs don't work as well for some problems:
 - unlikely to sample nodes in narrow passages
 - hard to sample/connect nodes on constraint surfaces

Literature

- J.C. Latombe, Robot Motion Planning, Kluwer Academic Publishers, 1991.
- S. LaValle, Planning Algorithms. 2006.
<http://msl.cs.uiuc.edu/planning/>
- Likhachev, ARA*, CMU
- Toussaint, Lecture Notes Robotics, 2011
- Dr. John (Jizhong) Xiao, City College New York

Literature, contd.

- (Limited) background in Russell&Norvig Chapter 25
- H. Choset et al., Principles of Robot Motion: Theory, Algorithms, and Implementations. 2006.
- Other demos/examples:
 - <http://voronoi.sbp.ri.cmu.edu/~choset/>
 - <http://www.kuffner.org/james/research.html>
 - <http://msl.cs.uiuc.edu/rrt/>